

Question 1

What level of SIMD/vector support does the CPU your computer provide?

List of SIMD provided by 4900HS:

1. AVX2, AVX
2. SSE, SSE3, SSSE3 SSE4.1, SSE4.2, SSE2

The SSE is just old type of AVX.

Questions 2, 3

What is the maximum operand size that your computer will support?

AVX2 provides the longest register length for data, in this case it's 256 bit.

Question 4

What is the clock speed of your CPU? (You may need to look this up via "About this Mac" or "lscpu").
Base: 3.0 GHz, Boost: 4.3 GHz, usually I get 3.9GHz.

Question 5

Based on the output from bandwidth.exe on your computer, what do you expect L1 cache and L2 cache sizes to be? What are the corresponding bandwidths? How do the cache sizes compare to what "about this mac" (or equivalent) tells you about your CPU? (There is no "right" answer for this question – but I do want you to do the experiment.)

This is the verbatim of the out of from "./bandwidth.exe" command:

```

      read
bytes/elt  #elts  res_bytes  ntrials      usecs      ttl_bytes      bytes/sec
8          16      128      67108868      15000      8589935104      5.72662e+11
8          32      256      33554436      8000       8589935616      1.07374e+12
8          64      512      16777220      4000       8589936640      2.14748e+12
8         128     1024      8388612      1000       8589938688      8.58994e+12
8         256     2048      4194308           0      8589942784           0
8         512     4096      2097156     1000      8589950976      8.58995e+12
8        1024     8192      1048580           0      8589967360           0
8        2048     16384      524292           0      8590000128           0
8        4096     32768      131074           0      4295032832           0
8        8192     65536      65538           0      4295098368           0
8       16384     131072      32770           0      4295229440           0
8       32768     262144      16386           0      4295491584           0
8       65536     524288       8194           0      4296015872           0
8      131072     1048576      2049           0      2148532224           0
8      262144     2097152      1025           0      2149580800           0
8      524288     4194304       513           0      2151677952           0
8     1048576     8388608       257           0      2155872256           0
8     2097152     16777216      129           0      2164260864           0
8     4194304     33554432       65           0      2181038080           0
8     8388608     67108864       33           0      2214592512           0
8    16777216    13421728      17           0      2281701376           0

      write
bytes/elt  #elts  res_bytes  ntrials      usecs      ttl_bytes      bytes/sec
8          16      128      67108868      62000      8589935104      1.38547e+11
8          32      256      33554436      62000      8589935616      1.38547e+11
8          64      512      16777220      61000      8589936640      1.40819e+11
8         128     1024      8388612      61000      8589938688      1.40819e+11
8         256     2048      4194308      61000      8589942784      1.40819e+11
8         512     4096      2097156      61000      8589950976      1.40819e+11
8        1024     8192      1048580      61000      8589967360      1.40819e+11
8        2048     16384      524292      61000      8590000128      1.4082e+11
8        4096     32768      131074      32000      4295032832      1.3422e+11
8        8192     65536      65538      33000      4295098368      1.30154e+11
8       16384     131072      32770      31000      4295229440      1.38556e+11
8       32768     262144      16386      31000      4295491584      1.38564e+11
8       65536     524288       8194      32000      4296015872      1.3425e+11
8      131072     1048576      2049      22000      2148532224      9.76606e+10

```

8	262144	2097152	1025	23000	2149580800	9.346e+10
8	524288	4194304	513	51000	2151677952	4.21898e+10
8	1048576	8388608	257	364000	2155872256	5.92273e+09
8	2097152	16777216	129	351000	2164260864	6.16599e+09
8	4194304	33554432	65	366000	2181038080	5.95912e+09
8	8388608	67108864	33	360000	2214592512	6.15165e+09
8	16777216	134217728	17	375000	2281701376	6.08454e+09
read/write						
bytes/elt	#elts	res_bytes	ntrials	usecs	ttl_bytes	bytes/sec
8	8	128	67108868	108000	8589935104	7.95364e+10
8	16	256	33554436	130000	8589935616	6.60764e+10
8	32	512	16777220	65000	8589936640	1.32153e+11
8	64	1024	8388612	44000	8589938688	1.95226e+11
8	128	2048	4194308	37000	8589942784	2.32161e+11
8	256	4096	2097156	34000	8589950976	2.52646e+11
8	512	8192	1048580	32000	8589967360	2.68436e+11
8	1024	16384	524292	31000	8590000128	2.77097e+11
8	2048	32768	131074	17000	4295032832	2.52649e+11
8	4096	65536	65538	32000	4295098368	1.34222e+11
8	8192	131072	32770	32000	4295229440	1.34226e+11
8	16384	262144	16386	33000	4295491584	1.30166e+11
8	32768	524288	8194	33000	4296015872	1.30182e+11
8	65536	1048576	2049	21000	2148532224	1.02311e+11
8	131072	2097152	1025	20000	2149580800	1.07479e+11
8	262144	4194304	513	44000	2151677952	4.89018e+10
8	524288	8388608	257	265000	2155872256	8.13537e+09
8	1048576	16777216	129	264000	2164260864	8.19796e+09
8	2097152	33554432	65	262000	2181038080	8.32457e+09
8	4194304	67108864	33	265000	2214592512	8.35695e+09
8	8388608	134217728	17	203000	2281701376	1.12399e+10

The compiler optimized out the read forloop, making the results invalid. I currently don't have a good idea on how to get the read speed yet.

For the write, a Plateau appeared between the block size of "65536" to "131072", indicating a threshold of 2^{16} bits for the L1 registers. Which is about "64kb" L1 Cache. The second Plateau appeared between size "524288" and "1048576" bytes, meaning that L2 could have a size of "512" kb. The third plateau happens around 8mb, and that could potentially be entering the L3 cache.

For the read/write bandwidth to register, it's a different story. It seems like, for blocks that are too small, it's unable to achieve optimal speed, and the speed drops out, around "256kb".

L1 Cache has speed around "1.3425e+11" bytes/s, which is about "134.25 gb/s". And the decrease L2 speed is about "90 gb/s", and for L3, it's about "6gb/s".

The specs of the Ryzen 4900HS has 512kb for L1 register, and 4.0mb for L2 Register, and I think L3 is shared between 2 cluster of cores so we are not worrying about that.

Therefore, it seems like there are 2 types of speed for L1 cache for this CPU, I mistaken the first Plateau as falling L1 to L2, but actually it's not.

I googled a bit, and here is the source for [4900HS cache organization](#).


For L1, there are 2 types of L1, L1L, L1D, and for L2, it's 512kb shared for one core while 4 ways stitched together for 4 coers.

L3 is 2 times 4 mb (8 mb in total) and between 2 groups of 4 cores.

Note: The Ryzen zen 2 CPU should also have a FMA4 (Fused, Multiply Add) SIMD instruction set for floating operations, however, it's not shown when "./cpuinfo583.exe" is executed.

Questions 6

Based on the output from bandwidth.exe on your computer, what do you expect L1 cache and L2 cache sizes to be? What are the corresponding bandwidths? How do the cache sizes compare to what "about this mac" (or equivalent) tells you about your CPU? (There is no "right" answer for this question – but I do want you to do the experiment.)



There are 3 drops in the data bandwidth from the experiment, at the size of 32kb, 512 kb, and 4mb. Each of them corresponds to part of the L1 cache (each core has independent 32kb of cache to work with), all of the L1 cache, L2. And by the time we entered L3 cache (around 4mb), all the speed is gone because we are in the DRAM territory.

Question 7

What is the (potential) maximum compute performance of your computer? (The horizontal line.) What are the L1, L2, and RAM bandwidths? How do those bandwidths correspond to what was measured above with the bandwidth program?

1. Assume single core performance, The top speed is computed by:

$$(\text{Core Speed}) \times (\text{Instruction per clcye}) \times (\text{Number of Floats it can process with SIMD})$$

hence the number will be given by:

$$4.3 \times 8[\text{GFlops/s}] \approx 34.4[\text{Gflops/s}]$$

where the data “8” is for double precision floating points, using the “FMA3” instruction for ryzen processor. However, if the “256” bits register is used with the “AVX2” SIMD, then the data Gflops will be $4.3 \times 4 \approx 17.2$. With both contributing together, we can have: $4.3 \times 8 + 4.3 \times 3 \approx 47.3$ Gflops per second.

I will say this is a very gross over estimate (It hardly even goes to the advertised 4.3 GHz to be honest.), and I will say it's about “15” Gflops under the best case for my CPU.

This part of the HW Questions is done with the help from [this](#) stack overflow discussion.

2. L1: “255 GB/s”, L2 is “88 GB/s”, RAM is “9 GB/s”, this is for single core I would assume.
3. My PC manufacture didn't specify any of this information, there is no comparison to make. The ram is slower then the expected value from the CPU, no idea why, but my machine is 8 + 32 gb config double channel. This is a bad idea.

Question 8

Based on the clock speed of your CPU and its maximum Glop rate, what is the (potential) maximum number of *double precision* floating point operations that can be done per clock cycle? (Hint: $\text{Glops} / \text{sec} : \text{math:} \div \text{GHz} = \text{flops} / \text{cycle}$.) There are several hardware capabilities that can contribute to supporting more than one operation per cycle: fused multiply add (FMA) and AVX registers. Assuming FMA contributes a factor of two, SSE contributes a factor of two, AVX/AVX2 contribute a factor of four, and AVX contributes a factor of eight of eight, what is the expected maximum number of floating point operations your CPU could perform per cycle, based on the capabilities your CPU advertises via `cpuid` (equiv. `lscpu`)? Would your answer change for single precision (would any of the previous assumptions change)?

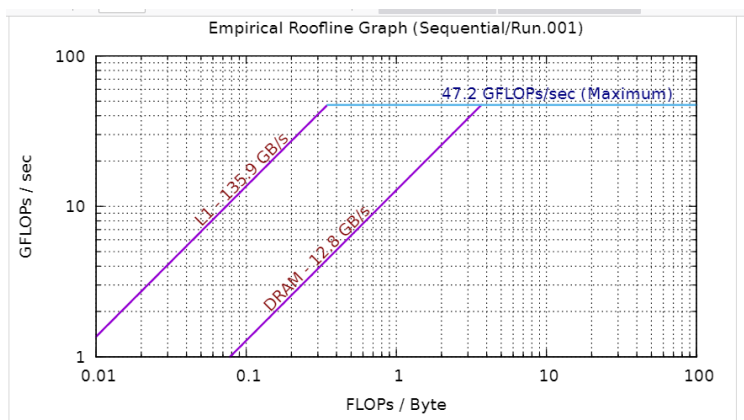
1. From retuls pulled from question, 9, which is 47.2 Gflops per second, then the Gflops per cycle is going to be: $47.2 / 4.3 \approx 10.9$ floats per second.

Theoretically speaking, the the Ryzen has “FMA3” and “AVX2”, the first can compute 8 double floating and 3 for FMA3, hence the theoretical flops per second is $8 + 3 = 11$.

2. My answer will change for single precision. The AVX, or the FMA instruction set can be divided up for streaming single precision numbers, this allows for the performance to double when computing single precision floating points compare to double precision floating points.

Question 9

What is the maximum compute performance of your computer? (The horizontal line.) What are the L1, L2, and DRAM bandwidths? How do those bandwidths correspond to what was measured above?



The maximal compute performance is 47.2 GFlops per second (Very closed to mixed AVX2 with FMA3 SIMDs). The L1 is 135.9 GB/s, the same as we observed from the homegrown benchmark. The L2 Cache bandwidth is gone under the eyes of this profiler or unknown reasons.

They are very similiar to the measurements above. The DRAM bandwidth is the same as results from the “amath583/bandwidth” docker run.

Back in the question 7, I assume the usage of both AVX2 with the FMA3 SIMDs, and gotten the maximal Flops throughput to be 51.6, which is the closest to what we had for the roofline in the graph computed via docker.


Question 10

Referring to the figures about how data are stored in memory, what is it about the best performing pair of loops that is so advantageous?

This is the verbatim of the out put from executing “mmult.exe”:

N	GF/s ijk	GF/s ikj	GF/s jik	GF/s jki	GF/s kij	GF/s kji
8	4.84614	4.15383	3.48922	2.72595	4.84614	3.11537
16	3.87691	4.81272	3.47186	3.00795	4.9846	3.92047
32	2.99043	4.97858	2.74576	3.87223	5.09045	4.67062
64	2.33768	5.03899	2.33167	2.55498	4.95638	2.83443
128	1.71237	4.71744	1.69354	1.49195	4.64628	1.44109
256	1.14627	4.61373	1.22828	0.39231	4.80389	0.398308

The ikj and kij ordering performs the best. Observe that, the the most inner loop is **both iterating on the index j** . The expression $C(i, j) += A(i, k) * B(k, j)$, where the fastest iterating index j is running over the columns of C , and columns of B . The key here is that, running over the same column in the inner loop increases the memory locality because the underlying storage of the matrix is row major, and it takes the form of “vector<double>”, and we indexed it using “[m*i + j]”. Notice that, incrementing on “j” will give us sequential access to the data. Increasing the memory locality of this numerical algorithm. Here is a graph, and we imagine how the indexing goes along the memory direction when the innermost loop is iterating through index “j” (stays on the same rows of the matrix C, B and goes from the top to the bottom)



Question 11


What will the data access pattern be when we are executing “mult_trans” in i,j,k order? What data are accessed in each if the matrices at step (i,j,k) and what data are accessed at step (i, j, k+1)? Are these accesses advantageous in any way?

The expression for multiplying AB^T is: $C(i, j) += A(i, k) * B(j, k)$. And take notice, the inner forloop controls index k , and it sweeps through the i th row of A and the j th row of B , from top to bottom, in a sequential manner.

This is very good because the fastest going loop is going through columns of A, B and the second fastest loop is going through the row of A, B , in addition, C is not indexed by the fastest index k , this maximizes the memory locality of this algorithm.

Question 12

Referring again to how data are stored in memory, explain why hoisting “C(i,j)” out of the inner loop is so beneficial in mult_trans with the “ijk” loop ordering.



The fastest running index “k” will sweep through several loops at the same time, and they will all be present with the same address space in the memory, take advantage of sequential read and increasing the speed of the cache and compiler optimization. While the slower running index move on a larger chunks of address space.

The hoisting allows the compiler to optimize and vectorized the accumulations of all the multiply adds.

Question 13

What optimization is applied in going from “mult_2” to “mult_3”?

Blocks of 32 by 32 is introduced in addition of tiling of 2 by 2 are used for each of the blocks.

This enables the compiler to infer more on optimizations for the cache, and the use of SMID. This also increase the memory locality because of the usage of 32 by 32 blocks.

Question 14

How does your maximum achieved performance for “mult” (any version) compare to what bandwidth and roofline predicted?

Question 15

Which variant of blurring function between struct of arrays and array of structs gives the better performance? Explain what about the data layout and access pattern would result in better performance.

Running on “julia.bmp”

SOA inner	SOA outer	AOS inner	AOS outer	Ten inner	Ten outer
31	5	5	14	31	4

Running on “med_julia.bmp”

SOA inner	SOA outer	AOS inner	AOS outer	Ten inner	Ten outer
505	83	86	268	574	96

Running on “large_julia.bmp”

SOA inner	SOA outer	AOS inner	AOS outer	Ten inner	Ten outer
2244	371	353	1086	9592	384

The performance of SOA outer and AOS inner is similar. The SOA outer has color as the outter most for loop, and in that case, filter of the same color, stored in vector are read. However, is slower because the memory each color are no sequentially pact together. The AOS inner mitigate the jumps between each color by packing array of size 3 together into one color. Hence, when interating 3 colors inside the forloop, the memory locality is maximized and the reading of each of the pixels are sequenal as we iterate through the vector of pixels.

SOA inner is definitely worse because it jumps between different colors of filter, and they are $N * N$ bytes apart from each other in the memory. Compare to AOS, where the it jumps between different pixels in the inner forloop, and the distance between them are 3 bytes.

Therefore, the speed of AOS is faster than SOA.

Question 16

Which variant of the blurring function has the best performance overall? Explain, taking into account not only the data layout and access pattern but also the accessor function.

AOS inner has the best performance. Because the inner loop jumps between different color, which are right next to each other in the array inside of the vector, and the outter loops jumps between different pixels, which is also pack sequentially in the vector.

Logs

mut_0.log

N	GF/s ijk	GF/s ikj	GF/s jik	GF/s jki	GF/s kij	GF/s kji
8	4.36152	3.79263	3.48922	2.02862	3.79263	2.29554
16	3.81335	4.74724	3.48922	3.06072	4.9846	3.89857
32	2.98059	4.84546	2.72102	3.83941	5.03389	4.44167
64	2.30793	4.82457	2.30208	2.5195	5.01115	2.8256
128	1.68691	4.66639	1.66477	1.4626	4.61643	1.44592
256	1.19385	4.75234	1.20752	0.39688	4.80389	0.393076

mult_1.log

N	GF/s ijk	GF/s ikj	GF/s jik	GF/s jki	GF/s kij	GF/s kji
8	7.93004	4.15383	3.48922	2.72595	4.84614	3.11537
16	15.8601	4.74724	3.37123	2.84834	4.68351	3.32307
32	24.4892	4.97858	2.72102	3.85575	5.06202	4.44167
64	19.2983	5.01115	2.32569	2.55498	5.01115	2.90711
128	10.9993	4.70714	1.67642	1.49299	4.71744	1.45667
256	6.0508	4.78314	1.22016	0.39688	4.81433	0.394828

mult_2.log

N	mult_0	mult_1	mult_2	mult_3	mul_t_0	mul_t_1	mul_t_2	mul_t_3
8	3.48922	3.6346	8.72305	9.69227	2.72595	5.4519	9.69227	14.5384
16	3.45467	3.57869	9.18215	18.3643	3.38759	4.10496	9.30458	25.8461
32	2.69673	2.77095	8.38982	25.8886	2.75411	3.02034	8.38982	32.3607
64	2.31974	2.5195	6.9238	19.7178	2.29625	2.91646	7.49602	20.156
128	1.71101	1.9581	7.04533	10.7794	1.9815	2.14089	6.63345	11.1127
256	1.34708	1.59323	6.23829	5.82788	1.84089	1.91078	6.9206	6.20334

mmult_ps3.log

N	mult_0	mult_1	mult_2	mult_3	mul_t_0	mul_t_1	mul_t_2	mul_t_3
8	3.48922	3.6346	8.72305	7.26921	4.84614	4.59108	9.69227	14.5384
16	3.50675	3.57869	9.18215	8.40776	3.87691	4.12925	9.30458	25.8461
32	2.77095	2.80527	8.54812	8.46823	2.95147	3.02034	8.38982	32.3607
64	2.31974	2.52651	6.87135	8.39832	2.33768	2.96411	7.622	21.0934
128	1.71373	1.98515	7.18624	7.926	2.0036	2.14301	7.23447	11.5287
256	1.33329	1.61885	6.69061	7.16697	1.86257	1.91243	7.0082	6.20334

Note: For consistency, all benchmark are run with the power of my laptop on, with maximal performance setting.