

Question 1

What level of SIMD/vector support does the CPU your computer provide?

List of SIMD provided by 4900HS:

1. AVX2, AVX
2. SSE, SSE3, SSSE3 SSE4.1, SSE4.2, SSE2

The SSE is just old type of AVX.

Questions 2, 3

What is the maximum operand size that your computer will support?

A in this case, AVX 2 provides the longest register length for data, in this case it's 256 bit.

Question 4

What is the clock speed of your CPU? (You may need to look this up via "About this Mac" or "lscpu").

Base: 3.0 GHz, Boost: 4.3 GHz, usually I get 3.9GHz.

Question 5

Based on the output from bandwidth.exe on your computer, what do you expect L1 cache and L2 cache sizes to be? What are the corresponding bandwidths? How do the cache sizes compare to what "about this mac" (or equivalent) tells you about your CPU? (There is no "right" answer for this question – but I do want you to do the experiment.)

This is the verbatim of the out of from "./bandwidth.exe" command:

```

read
bytes/elt      #elts  res_bytes  ntrials      usecs      ttl_bytes      bytes/sec
8              16      128        67108868     15000      8589935104     5.72662e+11
8              32      256        33554436     8000       8589935616     1.07374e+12
8              64      512        16777220     4000       8589936640     2.14748e+12
8              128     1024       8388612      1000       8589938688     8.58994e+12
8              256     2048       4194308      0          8589942784     0
8              512     4096       2097156      1000       8589950976     8.58995e+12
8              1024    8192       1048580      0          8589967360     0
8              2048    16384      524292       0          8590000128     0
8              4096    32768      131074       0          4295032832     0
8              8192    65536      65538        0          4295098368     0
8              16384   131072     32770        0          4295229440     0
8              32768   262144     16386        0          4295491584     0
8              65536   524288     8194         0          4296015872     0
8              131072  1048576    2049         0          2148532224     0
8              262144  2097152    1025         0          2149580800     0
8              524288  4194304    513          0          2151677952     0
8              1048576 8388608     257          0          2155872256     0
8              2097152 16777216    129          0          2164260864     0
8              4194304 33554432    65           0          2181038080     0
8              8388608 67108864    33           0          2214592512     0
8              16777216 134217728   17           0          2281701376     0

write
bytes/elt      #elts  res_bytes  ntrials      usecs      ttl_bytes      bytes/sec
8              16      128        67108868     62000      8589935104     1.38547e+11
8              32      256        33554436     62000      8589935616     1.38547e+11
8              64      512        16777220     61000      8589936640     1.40819e+11
8              128     1024       8388612      61000      8589938688     1.40819e+11
8              256     2048       4194308      61000      8589942784     1.40819e+11
8              512     4096       2097156      61000      8589950976     1.40819e+11
8              1024    8192       1048580      61000      8589967360     1.40819e+11
8              2048    16384      524292       61000      8590000128     1.4082e+11
8              4096    32768      131074       32000      4295032832     1.3422e+11
8              8192    65536      65538        33000      4295098368     1.30154e+11
8              16384   131072     32770        31000      4295229440     1.38556e+11
8              32768   262144     16386        31000      4295491584     1.38564e+11
8              65536   524288     8194         32000      4296015872     1.3425e+11
8              131072  1048576    2049         22000      2148532224     9.76606e+10
8              262144  2097152    1025         23000      2149580800     9.346e+10
8              524288  4194304    513          51000      2151677952     4.21898e+10
8              1048576 8388608     257          364000     2155872256     5.92273e+09
8              2097152 16777216    129          351000     2164260864     6.16599e+09
8              4194304 33554432    65           366000     2181038080     5.95912e+09
8              8388608 67108864    33           360000     2214592512     6.15165e+09
8              16777216 134217728   17           375000     2281701376     6.08454e+09

read/write
bytes/elt      #elts  res_bytes  ntrials      usecs      ttl_bytes      bytes/sec
8              8        128        67108868     108000     8589935104     7.95364e+10
8              16      256        33554436     130000     8589935616     6.60764e+10
8              32      512        16777220     65000      8589936640     1.32153e+11
8              64      1024       8388612      44000      8589938688     1.95226e+11
8              128     2048       4194308      37000      8589942784     2.32161e+11
8              256     4096       2097156      34000      8589950976     2.52646e+11
8              512     8192       1048580      32000      8589967360     2.68436e+11
8              1024    16384      524292       31000      8590000128     2.77097e+11
8              2048    32768      131074       17000      4295032832     2.52649e+11
8              4096    65536      65538        32000      4295098368     1.34222e+11

```

8	8192	131072	32770	32000	4295229440	1.34226e+11
8	16384	262144	16386	33000	4295491584	1.30166e+11
8	32768	524288	8194	33000	4296015872	1.30182e+11
8	65536	1048576	2049	21000	2148532224	1.02311e+11
8	131072	2097152	1025	20000	2149580800	1.07479e+11
8	262144	4194304	513	44000	2151677952	4.89018e+10
8	524288	8388608	257	265000	2155872256	8.13537e+09
8	1048576	16777216	129	264000	2164260864	8.19796e+09
8	2097152	33554432	65	262000	2181038080	8.32457e+09
8	4194304	67108864	33	265000	2214592512	8.35695e+09
8	8388608	134217728	17	203000	2281701376	1.12399e+10

The compiler optimized out the read forloop, making the results invalid. I currently don't have a good idea on how to get the read speed yet.

For the write, a Plateau appeared between the block size of "65536" to "131072", indicating a threshold of 2^{16} bytes for the L1 registers. Which is about "64kb" L1 Cache. The second Plateau appeared between size "524288" and "1048576" bytes, meaning that L2 could have a size of "512"kb. The third plateau happens around 8mb, and that could potentially be entering the L3 cache.

For the read/write bandwidth to register, it's a different story. It seems like, for blocks that are too small, it's unable to achieve optimal speed, and the speed drops out, around "256kb".

L1 Cache has speed around "1.3425e+11" bytes/s, which is about "134.25 gb/s". And the decrease L2 speed is about "90 gb/s", and for L3, it's about "6gb/s".

The specs of the Ryzen 4900HS has 512kb for L1 register, and 4.0mb for L2 Register, and I think L3 is shared between 2 cluster of cores so we are not worrying about that.

Therefore, it seems like there are 2 types of speed for L1 cache for this CPU, I mistaken the first Plateau as falling L1 to L2, but actually it's not.

I googled a bit, and here is the source for [4900HS cache organization](#).


For L1, there are 2 types of L1, L1L, L1D, and for L2, it's 512kb shared for one core while 4 ways stitched together for 4 coers.

L3 is 2 times 4 mb (8 mb in total) and shared for all 8 cores.

Note: The Ryzen zen 2 CPU should also have a FMA4 (Fused, Multiply Add) SIMD instruction set for floating operations, however, it's not shown when "./cpuinfo583.exe" is executed.

Questions 6

Based on the output from bandwidth.exe on your computer, what do you expect L1 cache and L2 cache sizes to be? What are the corresponding bandwidths? How do the cache sizes compare to what "about this mac" (or equivalent) tells you about your CPU? (There is no "right" answer for this question – but I do want you to do the experiment.)



There are 3 drops in the data bandwidth from the experiment, at the size of 32kb, 512 kb, and 4mb. Each of them corresponds to part of the L1 cache, all of the L1 cache, L2. And by the time we entered L3 cache (around 4mb), all the speed is gone.

Question 7

What is the (potential) maximum compute performance of your computer? (The horizontal line.) What are the L1, L2, and RAM bandwidths? How do those bandwidths correspond to what was measured above with the bandwidth program?

1. Assume single core performance, The top speed is computed by:

$$(\text{Core Speed}) \times (\text{Instruction per cycle}) \times (\text{Number of Floats it can process with SIMD})$$

hence the number will be given by:

$$4.3 \times 8 [\text{GFlops/s}] \approx 34.4 [\text{Gflops/s}]$$

where the data “8” is for double precision floating points, using the “FMA4” instruction for ryzen processor. However, if the “256” bits register is used with the “AVX2” SIMD, then the data Gflops will be $4.3 \times 4 \approx 17.2$. With both contributing together, we can have: $4.3 \times 8 + 4.3 \times 4 \approx 51.6$ Gflops per second.

I will say this is a very gross over estimate (It hardly even goes to the advertised 4.3 GHz to be honest.), and I will say it's about “15” Gflops under the best case for my CPU.

This part of the HW Questions is done with the help from [this](#) stack overflow discussion.

2. L1: “255 GB/s”, L2 is “88 GB/s”, RAM is “9 GB/s”, this is for single core I would assume.
3. My PC manufacture didn't specify any of this information, there is no comparison to make.

Question 8


Based on the clock speed of your CPU and its maximum Glop rate, what is the (potential) maximum number of *double precision* floating point operations that can be done per clock cycle? (Hint: Glops / sec

$\text{GHz} = \text{flops} / \text{cycle}$.) There are several hardware capabilities that can contribute to supporting more than one operation per cycle: fused multiply add (FMA) and AVX registers. Assuming FMA contributes a factor of two, SSE contributes a factor of two, AVX/AVX2 contribute a factor of four, and AVX contributes a factor of eight of eight, what is the expected maximum number of floating point operations your CPU could perform per cycle, based on the capabilities your CPU advertises via `cpuid` (equiv. `lscpu`)? Would your answer change for single precision (would any of the previous assumptions change)?

1. Attempts on computing the maximal performance of the CPU is already been shown in the first part of question 7.
2. My answer will change for single precision. The AVX, or the FMA instruction set can be divided up for streaming single precision numbers, this allows for the performance to double when computing single precision floating points compare to double precision floating points.

Question 9

What is the maximum compute performance of your computer? (The horizontal line.) What are the L1, L2, and DRAM bandwidths? How do those bandwidths correspond to what was measured above?



The maximal compute performance is 47.2 GFlops per second. The L1 is 135.9 GB/s, the same as we observed from the homegrown benchmark. The L2 Cache bandwidth is gone under the eyes of this profiler or unknown reasons.

They are very similar to the measurements above. The DRAM bandwidth is the same as results from the "amath583/bandwidth" docker run.